

In [13]:

```
#Importing Data
import warnings
warnings.filterwarnings('ignore')
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
df = pd.read_csv('musk_csv.csv')
```

In [14]:

```
print(df.shape)
df.head(5)
```

(6598, 170)

Out[14]:

	ID	molecule_name	conformation_name	f1	f2	f3	f4	f5	f6	f7	...	f158	f159
0	1	MUSK-211	211_1+1	46	-108	-60	-69	-117	49	38	...	-308	-309
1	2	MUSK-211	211_1+10	41	-188	-145	22	-117	-6	57	...	-59	-60
2	3	MUSK-211	211_1+11	46	-194	-145	28	-117	73	57	...	-134	-135
3	4	MUSK-211	211_1+12	41	-188	-145	22	-117	-7	57	...	-60	-61
4	5	MUSK-211	211_1+13	41	-188	-145	22	-117	-7	57	...	-60	-61

5 rows × 170 columns

In [15]:

```
print(df.molecule_name.describe())
print(df.conformation_name.describe())
```

```
count          6598
unique           102
top      NON-MUSK-j146
freq           1044
Name: molecule_name, dtype: object
count          6598
unique          6598
top      j146_4+5
freq            1
Name: conformation_name, dtype: object
```

Processing molecule_name

molecule_name has 102 unique values but the name also itself defines whether the chemical is MUSK or NON-MUSK i.e the classification task.

In [18]:

```
#Seperating features and classes
x = df.drop(['conformation_name', 'class', 'molecule_name', 'ID','molecule_name_processed'], axis = 1)
y = df['class']

#Converting classes vector (integers) to binary class matrix to use it with categorical_crossentropy.
from keras.utils import to_categorical
y = to_categorical(y)

#Train Test Split
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, Y_test = train_test_split(x, y, test_size = 0.2, random_state = 0)
print("Shape of Training Set = ", x_train.shape)
```

Shape of Training Set = (5278, 166)

Standardizing the features

In [19]:

```
from sklearn.preprocessing import StandardScaler
ss = StandardScaler()
x_train = ss.fit_transform(x_train)
x_test = ss.transform(x_test)
```

MLP Model

1. MLP Model is chosen as its simple, robust and works well for classification tasks.
2. As the data available for training is very less so deeper NN might not work well.
3. Dense layers with 64,32,32 units are used along with relu activation function.
4. To prevent overfitting l2 regularization and Dropout are used.
5. Categorical_crossentropy is used over binary_crossentropy to easily convert probabilistic output to class labels using argmax.

In [21]:

```
from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
from keras import regularizers
import keras

reg = regularizers.l2(0.001)
input_dim = x_train.shape[1]
batch_size = 64
n_epoch = 25
output_dim = 2

model = Sequential()

model.add(Dense(64, activation= 'relu',input_shape=(input_dim,), activity_regularizer=reg))
model.add(Dropout(0.20))
model.add(Dense(32, activation= 'relu',activity_regularizer = reg))
model.add(Dense(32, activation= 'relu', activity_regularizer = reg))
model.add(Dense(output_dim, activation = "sigmoid"))
model.summary()
model.compile(optimizer = "adam", loss='categorical_crossentropy', metrics=['accuracy'
])

result = model.fit(x_train,y_train, batch_size=batch_size, epochs=n_epoch, verbose=2, validation_data=(x_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_9 (Dense)	(None, 64)	10688
dropout_3 (Dropout)	(None, 64)	0
dense_10 (Dense)	(None, 32)	2080
dense_11 (Dense)	(None, 32)	1056
dense_12 (Dense)	(None, 2)	66
Total params: 13,890		
Trainable params: 13,890		
Non-trainable params: 0		

Train on 5278 samples, validate on 1320 samples

Epoch 1/25

- 1s - loss: 1.8836 - acc: 0.7779 - val_loss: 0.9234 - val_acc: 0.8485

Epoch 2/25

- 0s - loss: 0.7528 - acc: 0.8452 - val_loss: 0.6135 - val_acc: 0.8485

Epoch 3/25

- 0s - loss: 0.5590 - acc: 0.8452 - val_loss: 0.4936 - val_acc: 0.8485

Epoch 4/25

- 0s - loss: 0.4567 - acc: 0.8463 - val_loss: 0.4020 - val_acc: 0.8568

Epoch 5/25

- 0s - loss: 0.3788 - acc: 0.9032 - val_loss: 0.3512 - val_acc: 0.9492

Epoch 6/25

- 0s - loss: 0.3324 - acc: 0.9416 - val_loss: 0.3048 - val_acc: 0.9538

Epoch 7/25

- 0s - loss: 0.3003 - acc: 0.9515 - val_loss: 0.2754 - val_acc: 0.9636

Epoch 8/25

- 0s - loss: 0.2701 - acc: 0.9600 - val_loss: 0.2519 - val_acc: 0.9598

Epoch 9/25

- 0s - loss: 0.2489 - acc: 0.9665 - val_loss: 0.2334 - val_acc: 0.9712

Epoch 10/25

- 0s - loss: 0.2371 - acc: 0.9714 - val_loss: 0.2241 - val_acc: 0.9644

Epoch 11/25

- 0s - loss: 0.2284 - acc: 0.9651 - val_loss: 0.2669 - val_acc: 0.9462

Epoch 12/25

- 0s - loss: 0.2044 - acc: 0.9740 - val_loss: 0.2162 - val_acc: 0.9614

Epoch 13/25

- 0s - loss: 0.1929 - acc: 0.9756 - val_loss: 0.1826 - val_acc: 0.9742

Epoch 14/25

- 0s - loss: 0.1883 - acc: 0.9754 - val_loss: 0.1970 - val_acc: 0.9682

Epoch 15/25

- 0s - loss: 0.1825 - acc: 0.9756 - val_loss: 0.1644 - val_acc: 0.9811

Epoch 16/25

- 0s - loss: 0.1734 - acc: 0.9790 - val_loss: 0.1875 - val_acc: 0.9712

Epoch 17/25

- 0s - loss: 0.1732 - acc: 0.9763 - val_loss: 0.1802 - val_acc: 0.9727

Epoch 18/25

- 0s - loss: 0.1544 - acc: 0.9831 - val_loss: 0.1503 - val_acc: 0.9833

Epoch 19/25

- 0s - loss: 0.1595 - acc: 0.9797 - val_loss: 0.1707 - val_acc: 0.9742

Epoch 20/25

- 0s - loss: 0.1618 - acc: 0.9782 - val_loss: 0.1662 - val_acc: 0.9780

Epoch 21/25

- 0s - loss: 0.1511 - acc: 0.9799 - val_loss: 0.1652 - val_acc: 0.9750

Epoch 22/25

```
- 0s - loss: 0.1539 - acc: 0.9786 - val_loss: 0.1794 - val_acc: 0.9689
Epoch 23/25
- 0s - loss: 0.1438 - acc: 0.9812 - val_loss: 0.1393 - val_acc: 0.9750
Epoch 24/25
- 0s - loss: 0.1375 - acc: 0.9812 - val_loss: 0.1387 - val_acc: 0.9803
Epoch 25/25
- 0s - loss: 0.1357 - acc: 0.9831 - val_loss: 0.1256 - val_acc: 0.9886
```

Plotting and printing Metrics

In [22]:

```
test_loss = result.history['val_loss']
train_loss = result.history['loss']
test_acc = result.history['val_acc']
train_acc = result.history['acc']

n_epoch_lst = list(range(1,n_epoch+1))


import matplotlib.pyplot as plt
fig = plt.figure(figsize=(18,8))

#Plotting Train vs Validation Loss
ax1 = plt.subplot(121)
ax1.plot(n_epoch_lst,test_loss, label='Validation_loss' )
ax1.plot(n_epoch_lst,train_loss, label='Training_loss')
plt.legend()
plt.title('Train vs Validation Loss')

#Plotting Train vs Validation Accuracy
ax2 = plt.subplot(122)
ax2.plot(n_epoch_lst,test_acc, label='Validation_Accuracy' )
ax2.plot(n_epoch_lst,train_acc, label='Training_Accuracy')
plt.legend()
plt.title('Train vs Validation Accuracy')
plt.plot()
```

Out[22]:

[]



In [23]:

```
#Printing Training and Validation - Accuracy & Loss
score = model.evaluate(x_test, Y_test, verbose = 0)
print('Validation Loss',score[0])
print('Validation Accuracy',score[1]*100)
print('***20)
print('Train loss',train_loss[-1])
print('Train Accuracy',train_acc[-1]*100)

#Predicting classes from validation set
Y_pred = model.predict(x_test)


#Converting Probabilistic prediction Output to Class Labels (Musk or Non-Musk)
Activities = {0:'Non-Musk',1:'Musk'}
y_pred = pd.Series([Activities[y] for y in np.argmax(Y_pred, axis=1)])
y_test = pd.Series([Activities[y] for y in np.argmax(Y_test, axis=1)])

#Confusion Matrix
from sklearn.metrics import confusion_matrix
cm=confusion_matrix(y_test, y_pred)
print(cm)
labels=['Musk','Non-Musk']
cm_df = pd.DataFrame(cm,index=labels, columns =labels)

#Confusion Matrix Heatmap
import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize=(8,8))
sns.heatmap(cm_df, annot=True, fmt='d', cmap='cool', linewidth=5)
plt.xlabel('Predicted values')
plt.ylabel('True values')
plt.xticks(rotation=90)
plt.title('Confusion Matrix')
plt.show()
```

Validation Loss 0.09127355534018892
Validation Accuracy 98.86363636363636

Train loss 0.13566809801745477
Train Accuracy 98.31375521030694
[[191 9]
 [6 1114]]



In [24]:

```
#Using classification report to print other metrics (F1-score, Precision, Recall)  
from sklearn import metrics  
classificationreport = metrics.classification_report(y_test, y_pred)  
print(classificationreport)
```

	precision	recall	f1-score	support
Musk	0.97	0.95	0.96	200
Non-Musk	0.99	0.99	0.99	1120
micro avg	0.99	0.99	0.99	1320
macro avg	0.98	0.97	0.98	1320
weighted avg	0.99	0.99	0.99	1320